

---

# Global Prompt Refinement with Non-Interfering Attention Masking for One-Shot Federated Learning

---

Zhuang Qi<sup>1</sup>, Yu Pan<sup>1</sup>, Lei Meng<sup>1,\*</sup>, Sijin Zhou<sup>2</sup>, Han Yu<sup>3</sup>, Xiaoxiao Li<sup>4,5</sup>, Xiangxu Meng<sup>1</sup>

<sup>1</sup>School of Software, Shandong University, China

<sup>2</sup>AIM Lab, Faculty of Engineering, Monash University, Clayton, VIC, Australia

<sup>3</sup>College of Computing and Data Science, Nanyang Technological University, Singapore

<sup>4</sup>Department of Electrical and Computer Engineering, University of British Columbia, Canada

<sup>5</sup>Vector Institute, Canada

`z_qi@mail.sdu.edu.cn`, `202100300334@mail.sdu.edu.cn`, `lmeng@sdu.edu.cn`,  
`sjzhou1995@gmail.com`, `han.yu@ntu.edu.sg`, `xiaoxiao.li@ece.ubc.ca`, `mxx@sdu.edu.cn`

## Abstract

Federated Prompt Learning (FPL) enables communication-efficient adaptation by tuning lightweight prompts on top of frozen pre-trained models. Existing FPL methods typically rely on global information, which is only available after the second training round, to facilitate collaboration among client models. Therefore, they are inherently dependent on multi-round communication to fully exhibit their strengths. Moreover, existing one-shot federated learning methods typically focus on fitting seen tasks, but lack cross-task generalization. To bridge this gap, we propose the Global Prompt Refinement with Non-Interfering Attention Masking (GPR-NIAM) method for one-shot FPL. The core idea is to design a masking mechanism that restricts excessive interaction between the original text embeddings and the learnable prompt embeddings. GPR-NIAM achieves this through the collaboration of two key modules. Firstly, the attention isolation module suppresses attention from the learnable prompt tokens to the original text tokens, and reweights the reverse attention which preserves generalization across tasks. Secondly, the cross-silo collaborative refinement module integrates decentralized visual knowledge into a unified base and calibrates the global prompt through multi-source cross-modal knowledge alignment, further mitigating the inconsistency caused by data heterogeneity. Extensive experiments conducted on ten benchmark datasets under two tasks show that GPR-NIAM outperforms eight state-of-the-art methods in both class-level and domain-level generalization.

## 1 Introduction

Federated learning (FL) enables collaborative modeling across data sources in a privacy-preserving manner, which has gained traction across multiple domains [1, 2, 3, 4, 5]. It learns global knowledge by aggregating model weights or gradients instead of moving data to a central entity [6, 7, 8, 9]. Recently, Federated Prompt Learning (FPL) has emerged to support parameter-efficient fine-tuning in this setting. It not only leverages pre-trained models (e.g., CLIP [10]) more effectively, but also reduces communication overhead between FL clients and the FL server by transmitting only learnable prompts [11, 12, 13, 14]. However, heterogeneous data distributions across clients often degrade model aggregation performance and slow down convergence [15, 16]. This is because significant differences in the update directions across local models can cause conflicts during aggregation

---

\*Corresponding author




Figure 1: An overview of the proposed GPR-NIAM method. It supports three interaction types: full attention (0), hard masking ( $-\infty$ ), and reweighting ( $\lambda$ ). Hard masking blocks learnable tokens  $T_P$  from affecting original text tokens  $T_T$ , while  $\lambda$ -weighted links enable partial influence from special tokens  $T_{EOS}$ . The upper triangular mask follows causal attention for autoregressive decoding.

[17, 18, 19, 20, 21]. As a result, existing FPL methods still require multi-round communications between clients and the server to reconcile inconsistencies arising from client-specific training.

To deal with data heterogeneity, early FPL solutions (e.g., PromptFolio [17] and DP-FPL [22]) typically focused on incorporating global prompt information to constrain the update direction of local prompts in order to enhance the consistency and stability of federated optimization [17, 23, 22, 24]. However, since global prompt information is typically available only after the second communication round, these methods inherently rely on multi-round communications. This poses a challenge in real-world scenarios with limited communication resources. To this end, one-shot federated learning (OSFL), which requires only a single round of client-server interaction, offers a promising direction for further exploration [25, 26, 27, 28, 29]. OSFL significantly reduces the communication overhead by eliminating iterative exchanges, making it attractive for real-world deployments with strict latency or bandwidth constraints [30, 31, 32]. However, directly applying existing OSFL methods to FPL settings is not suitable, as they are typically optimized for seen tasks while neglecting the preservation of generalization to unseen tasks.

To bridge this important gap, we propose the Global Prompt Refinement with Non-Interfering Attention Masking (GPR-NIAM) method. It prevents the modification of the original token embeddings through a masking mechanism, thereby preserving transferable knowledge across tasks. As shown in Figure 1, the interaction between the learnable prompt embeddings and the original text token embeddings is unidirectional, which ensures that the latter remain unchanged. Specifically, GPR-NIAM consists of two main modules. Firstly, the Attention Isolation module selectively regulates the flow of prompt information through a novel masking strategy. It suppresses attention from the learnable prompt tokens to the original text tokens across all encoder layers, thereby effectively preventing alterations to the textual input token embeddings. In addition, it re-weights the reverse attention to encourage the model to extract task-relevant features more effectively from the original texts. Secondly, the cross-silo collaborative refinement (CSCR) module leverages visual representations extracted from multiple clients and utilizes multi-source visual supervision to refine the global prompt in a centralized manner. This further mitigates the negative impact of data heterogeneity.

We conduct extensive experiments on 10 datasets under two tasks: 1) base-to-base/novel generalization and 2) leave-one-domain-out generalization, and carry out performance comparisons, ablation studies, in-depth analysis and case studies. The results demonstrate that GPR-NIAM achieves superior overall performance by maintaining a strong balance between fitting seen tasks and generalizing to unseen ones, significantly outperforming eight competitive state-of-the-art approaches.

## 2 Related Work

### 2.1 Federated Prompt Learning

FPL aims to integrate FL with prompt learning by freezing the backbone model within the FL framework and introducing a set of learnable prompt parameters [17, 23, 22, 24, 33, 34, 35, 36]. This enables efficient personalization for on-device data while significantly reducing communication overhead, as only prompt parameters are shared between clients and the server. To mitigate the challenges posed by data heterogeneity across clients, popular FPL methods (e.g., FedOTP [24], FedPGP [23], PromptFolio [17], DP-FPL [22]) mostly adopt a shared prompt paradigm. In this design, a global prompt is constructed on the server and combined with client-specific information to

improve training stability and generalization under non-IID distributions. However, these methods are essentially designed for multi-round communications and require access to global feedback, which limits their effectiveness in one-shot FL scenarios.

## 2.2 One-Shot Federated Learning

OSFL is a communication-efficient variant of federated learning, where model training is completed within a single round of client-server interaction [25, 26, 37, 27, 28, 38, 39, 40, 41, 42]. Recent studies on OSFL mainly focus on generative and ensemble-based approaches. The former employs generative models to synthesize proxy data or representations on the server. For example, FedCVAE [41] uses a variational autoencoder to compress local knowledge. FedDISC [28] and FedDEO [38] leverage diffusion models to generate synthetic training data. The latter is dedicated to refining the fusion of model predictions from different clients. For example, FedOV [27] introduces an open-set recognition-based voting mechanism to address label skew in OSFL; IntactOFL [26] learns weighted or expert-based combinations to enhance overall performance. However, these methods are only designed for conventional tasks, aiming to fit the training objective while overlooking the preservation of cross-task generalization.

## 3 Preliminaries

Suppose there are  $K$  clients denoted by  $\mathcal{C} = \{C_1, C_2, \dots, C_K\}$ , and each client  $C_k \in \mathcal{C}$  holds a private dataset  $\mathcal{D}_k$ . The goal of traditional FL is to collaboratively train a global model  $f$  by minimizing the objective:  $\min_f \sum_{k=1}^K p_k \cdot \mathcal{L}(f, \mathcal{D}_k)$ , where  $p_k = \frac{|\mathcal{D}_k|}{\sum_{j=1}^K |\mathcal{D}_j|}$  is the weight of client  $k$ , and  $\mathcal{L}(\cdot)$  is a task-specific loss function [43, 44, 45, 46]. In FPL, each client  $k$  optimizes a set of learnable prompts  $\delta_k$  using its private dataset  $\mathcal{D}_k$ , while keeping the backbone model  $f$  frozen:  $\min_{\delta_k} \mathcal{L}(f; \delta_k, \mathcal{D}_k)$ . To support class-level prediction, the prompt for class  $c$  is constructed by concatenating the learnable prompt and a fixed class token. Specifically, let  $\mathbf{p}_c \in \mathbb{R}^{N_t \times d_t}$  denote the class token for class  $c$ , and  $\delta_k \in \mathbb{R}^{N_p \times d_t}$  denote the learnable prompt from client  $k$ , where  $N_t$  and  $N_p$  are the lengths of the fixed and learnable prompts,  $d_t$  is the dimension of hidden embedding, respectively. The full prompt can be formed as  $P_c^{(k)} = [\delta_k; \mathbf{p}_c]$ . The text feature for class  $c$  is computed by feeding the prompt into a text encoder  $\mathcal{T}(\cdot)$ :  $\mathbf{f}_t^c = \mathcal{T}(P_c^{(k)})$ . Given an input image  $\mathbf{x}$ , the image feature is extracted by the frozen image encoder  $f$ :  $\mathbf{f}_v = f(\mathbf{x})$ . The similarity score between the image and class- $c$  textual features is calculated as  $\rho_c = \text{Cosine}(\mathbf{f}_t^c, \mathbf{f}_v)$ , and the prediction probability is computed via softmax over similarities:  $p(\hat{y} = c \mid \mathbf{x}) = \frac{\exp(\rho_c/\tau)}{\sum_{j=1}^K \exp(\rho_j/\tau)}$ . The training objective is the cross-entropy loss between the similarity-based prediction and the ground-truth label:  $\mathcal{L}_{ce} = \ell(\boldsymbol{\rho}, \mathbf{e}_y)$ , where  $\boldsymbol{\rho} = [\rho_1, \dots, \rho_K]$ ,  $\mathbf{e}_y$  is the one-hot label vector, and  $\ell$  denotes the cross-entropy loss function. After local training, client  $C_k$  uploads its learned prompt  $\delta_k$  to the server. The server then aggregates the local prompts  $\{\delta_k\}_{k=1}^K$  to construct a global prompt  $\delta_g$ , typically by weighted averaging:  $\delta_g = \sum_{k=1}^K p_k \cdot \delta_k$ .

## 4 The Proposed GPR-NIAM Method

In this section, we propose the non-interfering attention masking method for one-shot federated prompt learning, called GPR-NIAM. Its core idea is to preserve the transferable knowledge by reducing the influence of learnable prompt tokens on the input text tokens, as illustrated in Figure 2.

### 4.1 The Attention Isolation (AttnIso) Module

Prompt tuning typically focuses on fitting local training tasks. However, data heterogeneity often leads to divergent optimization across clients and weakens its cross-task generalization. To address this, the AttnIso module designs attention masks to selectively block the information flow from learnable prompt tokens to original text tokens (e.g. “a photo of a [CLASS]”). We denote the text attention matrix as  $M \in \mathbb{R}^{(1+N_p+N_t)*(1+N_p+N_t)}$ , where  $N_p$  is the number of learnable prompts,  $N_t$  is the length of input textual feature tokens. The additional 1 corresponds to the [EOS] token. In addition, we apply customized attention control for the [EOS] token to ensure proper aggregation of contextual information from both the prompt and the input tokens.




Figure 2: Illustration of the GPR-NIAM framework. The Attention Isolation module preserves the original text embeddings to reduce the forgetting of generalized knowledge. The global prompt is refined using visual prototypes from all clients to mitigate data heterogeneity.

To achieve the above goal, the text attention matrix  $M$  can be formulated as follows,

$$M^{i,j} = \begin{cases} & \text{if } i \in [1, N_p] \wedge j \in [N_p + 1, N_p + N_t] \\ -\infty, & \text{or } i \in [N_p + 1, N_p + N_t] \wedge j \in [1, N_p] \quad (\text{Hard Masking}) \\ & \text{or } j = N_p + N_t + 1 \\ \lambda, & \text{if } i = N_p + N_t + 1 \wedge j \in [N_p + 1, N_p + N_t] \quad (\text{Reweighting}) \\ 0, & \text{otherwise} \quad (\text{Full Attention}) \end{cases} \quad (1)$$

where  $M^{i,j}$  denotes the  $i$ -th row and  $j$ -th column element of  $M$ , and  $\wedge$  denotes logical AND. Moreover, to integrate the attention mask  $M$  into the text encoder  $\mathcal{T}$ , we replace the original computation with the following:

$$\mathbf{y}^{(l+1)} = \mathcal{T}_{l+1}(\mathbf{y}^{(l)}, M) = \text{Concat}(h_1, h_2, \dots, h_H)W^O = [\tilde{p}_{l+1}^1, \dots, \tilde{p}_{l+1}^{N_p}, \tilde{t}_{l+1}^1, \dots, \tilde{t}_{l+1}^{N_t}, \tilde{s}_{l+1}^{\text{eos}}], \quad (2)$$

where  $h_i = \text{softmax}\left(\frac{Q_i K_i^T}{\sqrt{d_k}} + M\right) V_i, \quad i = 1, \dots, H$

where  $\mathbf{y}^{(l)} \in \mathbb{R}^{(N_p+N_t+1) \times d_t}$  and  $\mathbf{y}^{(l+1)} \in \mathbb{R}^{(N_p+N_t+1) \times d_t}$  are the input and output of the  $(l+1)$ -th encoder layer.  $\tilde{p}_{l+1}^i$ ,  $\tilde{t}_{l+1}^j$ , and  $\tilde{s}_{l+1}^{\text{eos}}$  represent the output embeddings of the  $i$ -th prompt token, the  $j$ -th text token, and the [EOS] token, respectively.  $Q_i, K_i, V_i \in \mathbb{R}^{(N_p+N_t+1) \times d_k}$  are the query, key, and value matrices for the  $i$ -th attention head,  $d_k = d_t/H$  denotes the hidden dimension of each attention head.  $\text{Concat}(h_1, \dots, h_H) \in \mathbb{R}^{(N_p+N_t+1) \times H d_k}$  denotes the concatenated outputs of all heads, and  $W^O \in \mathbb{R}^{H d_k \times d_t}$  is the output projection matrix mapping the result back to  $d_t$ . Following the masked encoding process, we obtain the final output of the text encoder as:  $\mathbf{y}^{(L)} = [\tilde{p}_L^1, \dots, \tilde{p}_L^{N_p}, \tilde{t}_L^1, \dots, \tilde{t}_L^{N_t}, \tilde{s}_L^{\text{eos}}]$ . Afterwards,  $\tilde{s}_L^{\text{eos}}$  is projected into a shared embedding space via a pre-trained head  $H_t(\cdot)$ . On the visual side, the input image  $\mathbf{x}$  is encoded by a frozen vision encoder  $f(\cdot)$  to obtain the visual representation. Overall, it focuses on minimizing the objective:

$$\min_{\delta_1, \dots, \delta_K} \sum_{k=1}^K \frac{|\mathcal{D}_k|}{|\mathcal{D}|} \cdot \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{D}_k} \left[ -\log \left( \frac{\exp(\text{sim}(f_v(\mathbf{x}), H_t(\tilde{s}_{L,k}^{(y)})))}{\sum_{c=1}^C \exp(\text{sim}(f_v(\mathbf{x}), H_t(\tilde{s}_{L,k}^{(c)})))} \right) \right], \quad (3)$$

where  $\tilde{s}_{L,k}^{(c)}$  denotes client  $k$ 's [EOS] embedding for class  $c$ ,  $\text{sim}(\cdot, \cdot)$  denotes cosine similarity function. The learnable prompts can be optimized by gradient descent in all clients:

$$\delta_k \leftarrow \delta_k - \eta \cdot \frac{1}{|\mathcal{B}_k|} \sum_{(\mathbf{x}, y) \in \mathcal{B}_k} \nabla_{\delta_k} \left[ -\log \left( \frac{\exp(\text{sim}(f_v(\mathbf{x}), H_t(\tilde{s}_{L,k}^{(y)})))}{\sum_{c=1}^C \exp(\text{sim}(f_v(\mathbf{x}), H_t(\tilde{s}_{L,k}^{(c)})))} \right) \right], \quad \forall k \in [K], \quad (4)$$

where  $\eta$  denotes the learning rate,  $\mathcal{B}_k$  is a batch data.

## 4.2 Cross-Silo Collaborative Refinement (CSCR) Module

To further alleviate the inconsistency caused by data heterogeneity, the CSCR module calibrates the global prompt through prototype-guided refinement across clients. To support this, each client  $C_k$

extracts class-wise visual prototypes from its local dataset  $\mathcal{D}_k$ . For class  $c$ , the prototype is computed as a randomly weighted average of the feature vectors from local samples in that class:

$$f_{v,k}^c = \sum_{\mathbf{x} \in \mathcal{D}_k^c} \alpha_{\mathbf{x}} \cdot f_v(\mathbf{x}), \quad \text{where } \sum_{\mathbf{x} \in \mathcal{D}_k^c} \alpha_{\mathbf{x}} = 1, \quad \alpha_{\mathbf{x}} \sim \mathcal{U}(0, 1). \quad (5)$$

where  $\alpha_{\mathbf{x}}$  denotes a random weight for each sample  $\mathbf{x}$  in class  $c$ , drawn from a uniform distribution and normalized to ensure the weights sum to 1. Notably, we can randomly sample the weight vector  $n$  times to generate multiple diverse prototypes per class  $\mathbf{P}_k^c = [f_{v,k}^{c,1}, f_{v,k}^{c,2}, \dots, f_{v,k}^{c,n}]$ , which enhances representational diversity, which differs from existing methods [47, 48, 49, 50, 51].  $f_{v,k}^{c,i}$  denotes  $i$ -th visual prototypes of class  $c$  in client  $k$ . Afterwards, each client  $C_k$  sends  $\mathbf{P}_k = \{\mathbf{P}_k^1, \dots, \mathbf{P}_k^N\}$  to the server, which aggregates all received prototypes into a global pool  $\mathcal{P}$ :

$$\mathcal{P} = \bigcup_{k=1}^K \mathbf{P}_k = \{\mathbf{P}_k^c \mid k \in [K], c \in [n]\}, \quad (6)$$

To calibrate the global prompt  $\delta_g$ , we leverage the global prototype pool  $\mathcal{P} = \{\mathcal{P}^1, \dots, \mathcal{P}^C\}$  and enforce the alignment between textual representations and visual prototypes via a class-level contrastive objective. Specifically, we design the following loss function:

$$\min_{\delta_g} -\frac{1}{|\mathcal{P}|} \sum_{f \in \mathcal{P}} \log \sigma \left( \text{sim}(f, H_t(\mathcal{T}(\delta_g, t_c))) - \log \sum_{j=1}^C \exp(\text{sim}(f, H_t(\mathcal{T}(\delta_g, t_j)))) \right), \quad (7)$$

where  $\sigma(\cdot)$  is the sigmoid function,  $t_c$  denotes 'a photo of c-th classname'. We then follow the gradient descent method to obtain the optimal parameters:

$$\delta_g \leftarrow \delta_g - \eta \cdot \frac{1}{|\mathcal{P}|} \sum_{(f,c) \in \mathcal{P}} \nabla_{\delta_g} \log \sigma \left( \text{sim}(f, H_t(\mathcal{T}(\delta_g, t_c))) - \log \sum_{j=1}^C \exp(\text{sim}(f, H_t(\mathcal{T}(\delta_g, t_j)))) \right), \quad (8)$$

where the global prompt  $\delta_g$  is initialized as the weighted average of all client-specific prompts (i.e.,  $\delta_g \leftarrow \sum_{k=1}^K \frac{|\mathcal{D}_k|}{|\mathcal{D}|} \cdot \delta_k$ ), and  $\eta$  is a learning rate.

### 4.3 Training Procedure

In this section, we provide the detailed training procedure of GPR-NIAM in Algorithm 1. As a two-stage optimization strategy, GPR-NIAM consists of a local prompt tuning phase and a global prompt refinement phase. Notably, it can be easily integrated into multiple prompt learning methods such as CoOp and TCP. In the first stage, each client initializes its local prompt from the global prompt and updates it on private data with masked attention learning. Simultaneously, class-wise visual prototypes are obtained by performing randomly weighted aggregation over local feature representations. Subsequently, the global prompt is optimized to align with all class prototypes in this pool via cross-silo collaborative retraining. With the two-stage design, GPR-NIAM alleviates the challenges posed by heterogeneous client data and promotes the retention of transferable knowledge for unseen tasks.

Moreover, the local prompt tuning stage is optional. Firstly, this stage provides a more meaningful initialization for the global prompt by adapting it to all clients data. Secondly, it generates class-wise visual prototypes to guide global refinement. While the framework can still perform global refinement without this phase, since it does not affect prototype modeling, omitting it might result in less effective initialization and slightly limit the overall refinement performance.

## 5 Experimental Evaluation

### 5.1 Experiment Settings

**Datasets** Following existing works [17, 52], experiments are conducted on CIFAR10 [53], Oxford-Pets [54], Caltech101 [55], DTD [56], FGVC Aircraft [57], Flowers102 [58], StanfordCars [59], and UCF101 [60] to evaluate base-to-base/novel generalization, and on Office-Home [61] and DomainNet [62] to assess leave-one-domain-out generalization. Their statistics are shown in Table 5.

---

**Algorithm 1** GPR-NIAM

---

```

1: Initialize the global prompt parameter  $\delta^0$ 
2: // Local Prompt Tuning Stage
3: for each client  $C_k \in \mathcal{C}$  in parallel do
4:   Initialize local prompt parameter  $\delta_k = \delta^0$ 
5:   for  $e = 1, \dots, E$  do
6:     Sample a mini-batch  $\zeta$  from local data  $\mathcal{D}_k$ 
7:     Update local prompt  $\delta_k$  using masked attention by Eq. (4)
8:   end for
9:   Compute visual prototype features for each class using Eq. (5)
10: end for
11: // Global Prompt Refinement Stage
12: Aggregate all visual prototypes to form a global prototype pool  $\mathcal{P}$  by Eq. (6)
13: for  $e' = 1, \dots, E'$  do
14:   Sample a mini-batch of prototypes  $\zeta'$  from  $\mathcal{P}$ 
15:   Update global prompt  $\delta_g$  via prototype-guided alignment using Eq. (8)
16: end for

```

---

**Evaluation Metrics** Following studies [47, 49, 63], we report the Top-1 Accuracy (Acc) results. For base-to-base/novel generalization tasks, we also compute the harmonic mean ( $HM = \frac{2 \cdot Acc_{base} \cdot Acc_{novel}}{Acc_{base} + Acc_{novel}}$ ) of accuracy on base ( $Acc_{base}$ ) and novel classes ( $Acc_{novel}$ ) to evaluate overall performance.

**Implementation Details** In all experiments, we set local training and global refinement epochs to 10, using the SGD optimizer with a learning rate of 0.001, a weight decay of 0.001, and a batch size of

Table 1: Performance comparison between GPR-NIAM and baselines on 8 datasets with  $\beta = 0.5$  in the base-to-base/novel generalization cases.

(a) Average over 8 datasets				(b) CIFAR10				(c) OxfordPets			
Method	Base	Novel	HM	Method	Base	Novel	HM	Method	Base	Novel	HM
CLIP	65.79	74.32	69.70	CLIP	95.30	95.50	95.39	CLIP	81.53	95.33	87.89
PromptFL	72.57	72.80	72.64	PromptFL	95.14	96.26	95.69	PromptFL	91.31	93.70	92.48
RPO	70.52	73.63	71.99	RPO	95.68	94.74	95.20	RPO	89.00	95.49	92.13
FedTPG	70.46	73.09	71.69	FedTPG	95.38	96.56	95.96	FedTPG	91.66	95.13	93.36
TCP	72.46	74.60	73.41	TCP	95.68	96.52	96.09	TCP	90.01	95.39	92.62
PromptFolio	73.25	73.35	73.15	PromptFolio	<b>96.04</b>	96.16	96.06	PromptFolio	91.13	<b>96.10</b>	93.54
DP-FPL	71.90	71.75	69.88	DP-FPL	95.42	95.18	95.29	DP-FPL	91.97	94.80	93.36
FedDISC	56.43	54.84	53.87	FedDISC	94.30	92.82	93.55	FedDISC	83.83	72.07	77.50
FedELMY	<b>76.62</b>	48.23	60.03	FedELMY	95.32	40.00	79.66	FedELMY	89.28	76.47	82.38
GPR-NIAM <sub>P</sub>	75.56	73.28	74.24	GPR-NIAM <sub>P</sub>	95.66	<b>97.00</b>	<b>96.32</b>	GPR-NIAM <sub>P</sub>	91.18	95.89	93.47
GPR-NIAM <sub>T</sub>	76.43	<b>74.48</b>	<b>75.30</b>	GPR-NIAM <sub>T</sub>	95.84	96.48	96.15	GPR-NIAM <sub>T</sub>	<b>92.48</b>	95.14	<b>93.79</b>
(d) Caltech101				(e) DTD				(f) FGVC Aircraft			
Method	Base	Novel	HM	Method	Base	Novel	HM	Method	Base	Novel	HM
CLIP	82.27	94.05	87.77	CLIP	53.35	<b>59.05</b>	56.06	CLIP	26.51	31.21	28.67
PromptFL	88.27	93.88	90.99	PromptFL	60.06	53.86	56.79	PromptFL	31.19	30.91	31.05
RPO	89.90	93.80	91.81	RPO	54.62	59.17	56.81	RPO	29.27	30.43	29.84
FedTPG	89.12	92.61	90.83	FedTPG	56.71	55.43	56.06	FedTPG	29.57	30.25	29.90
TCP	88.20	<b>95.33</b>	91.62	TCP	61.68	55.48	58.41	TCP	30.77	<b>32.71</b>	31.71
PromptFolio	87.28	93.12	90.10	PromptFolio	61.11	54.58	57.66	PromptFolio	31.37	32.39	<b>31.87</b>
DP-FPL	87.92	93.46	90.60	DP-FPL	62.26	49.63	55.24	DP-FPL	28.85	28.15	28.49
FedDISC	44.75	85.22	58.68	FedDISC	<b>71.52</b>	38.28	49.87	FedDISC	20.87	13.20	16.17
FedELMY	87.71	89.21	88.45	FedELMY	70.13	16.18	26.29	FedELMY	<b>33.77</b>	19.48	24.70
GPR-NIAM <sub>P</sub>	91.59	94.22	92.89	GPR-NIAM <sub>P</sub>	68.40	53.50	60.04	GPR-NIAM <sub>P</sub>	30.95	29.41	30.16
GPR-NIAM <sub>T</sub>	<b>92.65</b>	94.73	<b>93.68</b>	GPR-NIAM <sub>T</sub>	69.90	55.91	<b>62.13</b>	GPR-NIAM <sub>T</sub>	32.15	31.12	31.62
(g) Flowers102				(h) StanfordCars				(i) UCF101			
Method	Base	Novel	HM	Method	Base	Novel	HM	Method	Base	Novel	HM
CLIP	62.39	78.10	69.36	CLIP	56.87	69.40	62.51	CLIP	68.12	71.92	69.97
PromptFL	79.88	75.15	77.44	PromptFL	62.96	67.30	65.02	PromptFL	71.82	71.55	71.68
RPO	74.92	76.21	75.56	RPO	58.64	68.73	63.29	RPO	72.19	70.49	71.33
FedTPG	72.30	77.26	74.70	FedTPG	58.67	69.78	63.74	FedTPG	70.34	67.74	69.01
TCP	77.25	<b>78.73</b>	77.98	TCP	60.19	71.18	65.22	TCP	75.93	71.52	73.65
PromptFolio	81.04	74.51	77.63	PromptFolio	61.22	<b>69.91</b>	65.27	PromptFolio	76.88	69.75	73.14
DP-FPL	74.63	77.26	75.92	DP-FPL	59.07	66.05	62.36	DP-FPL	75.14	69.49	72.20
FedDISC	24.95	32.21	28.11	FedDISC	42.77	49.92	46.07	FedDISC	68.49	55.03	61.03
FedELMY	<b>89.79</b>	55.15	68.33	FedELMY	<b>67.09</b>	47.97	55.94	FedELMY	<b>79.94</b>	41.41	54.56
GPR-NIAM <sub>P</sub>	84.54	78.52	81.42	GPR-NIAM <sub>P</sub>	63.59	67.48	65.48	GPR-NIAM <sub>P</sub>	78.57	70.28	74.19
GPR-NIAM <sub>T</sub>	86.88	77.05	<b>81.67</b>	GPR-NIAM <sub>T</sub>	63.54	69.18	<b>66.24</b>	GPR-NIAM <sub>T</sub>	78.04	<b>76.27</b>	<b>77.14</b>

Table 2: Performance comparison on Office-Home and DomainNet under the leave-one-domain-out setting, with each domain’s data randomly split across 3 clients.

Method	Office-Home					DomainNet						
	Art	Clipart	Product	RealWorld	Average	Clipart	Infograph	Painting	Quickdraw	Real	Sketch	Average
CLIP	74.57	62.33	80.17	79.38	74.08	92.87	77.02	85.35	43.06	98.61	94.22	81.85
PromptFL	74.24	66.20	82.47	79.77	75.17	93.87	80.82	87.76	49.47	98.72	94.82	84.24
RPO	74.86	64.33	84.31	80.74	76.06	93.53	78.32	86.15	49.35	98.83	94.87	83.50
FedTPG	74.04	64.55	83.32	80.01	75.48	<b>95.14</b>	80.93	88.25	52.89	98.97	95.40	85.26
TCP	75.31	66.20	83.17	80.97	76.41	94.13	79.80	88.05	45.85	98.80	95.34	83.66
PromptFolio	77.21	67.74	84.74	81.98	77.92	94.87	<b>81.19</b>	88.61	47.59	98.85	95.02	84.35
DP-FPL	75.03	65.91	83.71	79.82	76.12	95.03	78.41	86.74	54.03	98.83	94.85	84.64
FedDISC	72.27	65.40	81.77	79.11	74.63	93.82	80.42	86.17	43.74	98.66	94.05	82.81
FedELMY	75.11	66.34	84.12	79.69	76.31	94.21	77.65	86.81	48.78	98.83	94.88	83.53
GPR-NIAM <sub>P</sub>	78.32	71.40	87.61	84.96	80.57	94.72	80.13	<b>88.70</b>	<b>56.98</b>	<b>99.18</b>	94.88	85.76
GPR-NIAM <sub>T</sub>	<b>79.39</b>	<b>72.06</b>	<b>88.03</b>	<b>85.49</b>	<b>81.24</b>	95.03	80.82	87.92	56.32	99.08	<b>95.60</b>	<b>85.79</b>

32. The number of communication rounds is 1. We simulate 50 clients for CIFAR-10 and 10 clients for other datasets in base-to-base/novel generalization, while using 9 clients for Office-Home and 15 for DomainNet. To simulate non-IID data, we adopt a Dirichlet distribution with parameter  $\beta = 0.5$ . Each client generates  $n \in \{5, 10, 20\}$  visual prototypes per class, and the reweighting parameter  $\lambda$  is selected from  $\{0.2, 0.5, 0.7, 1.0\}$ . More details and results can be found in the Appendix. And each client has one NVIDIA RTX 3090 with 24 GB GPU for training.

## 5.2 Results and Discussion

This section compares GPR-NIAM-enhanced variants (GPR-NIAM<sub>P</sub> and GPR-NIAM<sub>T</sub>, derived by integrating GPR-NIAM into PromptFL and TCP, respectively) with nine baselines: Zero-Shot CLIP [10] with hand-crafted text prompt template, e.g., “a photo of a [class]” [64], PromptFL [65], RPO [66], FedTPG [52], TCP [67], PromptFolio [17], DP-FPL [22], FedDISC [28], FedELMY [42]. For fair comparison, all methods use the same backbone: a frozen CLIP with ViT-B16 [10]. Following prior work, the learnable prompt is a  $10 \times 512$  parameter matrix, with 10 prompt tokens and a 512-dimensional embedding.

**Base-to-Base/Novel Generalization.** We evaluate the generalization capability of GPR-NIAM by comparing its performance on base and novel classes, and the overall harmonic mean (HM) across all benchmarks. The results are presented in Table 1 and key observations are summarized as follows:

- GPR-NIAM attains second-best performance on base classes among all methods, showing particular competitiveness against FedMLEY while maintaining more balanced capabilities across all classes.
- GPR-NIAM yields competitive performance on novel classes. Notably, it ranks first on Flowers102 and UCF101, underscoring its strong generalization capability to unseen tasks. Moreover, FedDISC may fail on both base and novel classes due to inevitable gaps between synthetic and real samples.
- Benefiting from the synergy of attention masking and cross-silo collaborative retraining, GPR-NIAM achieves the best HM scores across all benchmarks.

**Leave-One-Domain-Out Generalization.** Table 2 presents the domain generalization results under a leave-one-domain-out setting. The main findings are summarized below:

- GPR-NIAM consistently achieve the highest average accuracy across most domains. Notably, GPR-NIAM<sub>T</sub> ranks first in 9 out of 12 domains, highlighting its strength in cross-domain generalization, especially in challenging cases like Infograph and Quickdraw.
- By combining token-level and client-level perspectives, GPR-NIAM outperforms single-view methods such as PromptFL and FedTPG, which benefits from attention masking for preserving original semantics and a unified global prompt for cross-domain knowledge integration.

## 5.3 Ablation Study

This section presents an ablation study to examine the effects of key components in GPR-NIAM, including the Attention Isolation (AttnIso) module and the Cross-Silo Collaborative Retraining (CSCR) module. The results are reported in Table 3.

Table 3: Ablation results averaged over 8 datasets (base-to-base/novel generalization) and across all domains (leave-one-domain-out generalization).

Method	Base to Base/Novel Generalization			Leave-One-Domain-Out Generalization	
	Base	Novel	HM	Office-Home	DomainNet
PromptFL	64.06	71.48	66.98	75.17	84.24
+ AttnIsoHA	62.19	<b>72.27</b>	67.04	77.26	84.87
+ AttnIsoRE	64.22	71.51	67.76	75.67	84.36
+ AttnIsoHA+RE	64.11	72.66	67.87	77.83	85.01
+ CSCR	<b>73.34</b>	64.89	67.24	78.59	85.11
+ AttnIsoHA + CSCR	69.32	70.78	70.06	79.13	85.46
+ AttnIsoRE + CSCR	70.56	68.11	69.65	78.42	85.23
+ AttnIsoHA+RE + CSCR	72.32	69.47	<b>70.72</b>	<b>80.57</b>	<b>85.76</b>

- Different configurations of the AttnIso module exhibit distinct strengths. Hard masking (AttnIso<sub>HA</sub>) effectively suppresses token interference and enhances novel-class generalization, while reweighting (AttnIso<sub>RE</sub>) provides a more stable improvement. Their combination leads to improved overall generalization, showing that the two strategies are complementary.
- The CSCR module plays a vital role in optimizing base-class performance by learning a unified prompt that mitigates inter-client discrepancies, which can alleviate the negative impact of data heterogeneity. However, its contribution to novel-class generalization remains limited due to insufficient adaptability to unseen categories.
- By integrating both modules, GPR-NIAM achieves an optimal balance between acquiring new knowledge and preserving generalization. Importantly, this also suggests a promising direction for future work by combining complementary strategies rather than relying on a single one.




Figure 3: The impact of hyperparameters on performance is evaluated by varying  $\lambda$  from  $\{0, 0.2, 0.5, 0.7, 1.0\}$ , the number of prototypes per class  $n$  from  $\{3, 5, 10\}$  on three datasets.

#### 5.4 Sensitivity Analysis

This section evaluates the robustness of GPR-NIAM across varying hyperparameter settings on the DTD, Flowers102 and Office-Home datasets. Specifically, we investigate the effects of hyperparameters  $n$  (the number of prototypes generated per class on the client side) and  $\lambda$  (the weighting factor in the AttnIso module), evaluated over  $\{3, 5, 10\}$  and  $\{0, 0.2, 0.5, 0.7, 1.0\}$ , respectively. As depicted in Figure 3, GPR-NIAM consistently achieves higher overall HM scores than its baselines in all cases of the base-to-base/novel generalization task, showcasing its superior balance between retaining transferable knowledge and learning task-specific knowledge. On the one hand, GPR-NIAM demonstrates consistent improvements over baselines on base classes, which benefits from the CSCR module’s capacity to alleviate data heterogeneity. On the other hand GPR-NIAM exhibits competitive performance on novel-class generalization, frequently outperforming other methods, which is attributed to the attention masking mechanism of the AttnIso module. Furthermore, in the Leave-One-Domain-Out Generalization task, similar conclusions are drawn, as GPR-NIAM outperforms the baseline across multiple domains.

Table 4: Performance comparison between GPR-NIAM and baselines on DTD and Flowers102 with local training epochs  $\in \{1, 5, 20\}$  in the base-to-base/novel generalization setting.

(a) DTD (Epochs=1)				(b) DTD (Epochs=5)				(c) DTD (Epochs=20)			
Method	Base	Novel	HM	Method	Base	Novel	HM	Method	Base	Novel	HM
CLIP	53.35	59.05	56.06	CLIP	53.35	59.05	56.06	CLIP	53.35	59.05	56.06
PromptFL	50.34	51.69	51.01	PromptFL	55.55	56.76	56.15	PromptFL	63.31	51.08	56.54
RPO	52.77	<b>59.42</b>	55.90	RPO	53.24	<b>59.78</b>	56.32	RPO	53.93	<b>60.26</b>	56.92
FedTPG	40.27	36.47	38.28	FedTPG	49.88	47.22	48.51	FedTPG	56.25	56.28	56.26
TCP	53.24	53.38	53.31	TCP	57.63	56.40	57.01	TCP	63.07	59.54	61.25
PromptFolio	52.54	53.50	53.02	PromptFolio	58.68	56.28	57.45	PromptFolio	58.68	57.60	58.13
DP-FPL	43.75	41.54	42.61	DP-FPL	59.25	50.12	54.30	DP-FPL	64.12	47.82	54.78
GPR-NIAM <sub>P</sub>	<b>65.97</b>	51.20	<b>57.65</b>	GPR-NIAM <sub>P</sub>	<b>68.51</b>	52.65	59.54	GPR-NIAM <sub>P</sub>	<b>70.13</b>	53.26	60.54
GPR-NIAM <sub>T</sub>	63.07	51.81	56.89	GPR-NIAM <sub>T</sub>	65.85	55.07	<b>59.98</b>	GPR-NIAM <sub>T</sub>	69.67	54.83	<b>61.36</b>
(d) Flowers102 (Epochs=1)				(e) Flowers102 (Epochs=5)				(f) Flowers102 (Epochs=20)			
Method	Base	Novel	HM	Method	Base	Novel	HM	Method	Base	Novel	HM
CLIP	62.39	78.10	69.36	CLIP	62.39	78.10	69.36	CLIP	62.39	78.10	69.36
PromptFL	70.84	74.10	72.43	PromptFL	78.42	76.00	77.19	PromptFL	80.46	73.68	76.92
RPO	69.97	76.21	72.95	RPO	71.72	76.84	74.19	RPO	74.05	77.05	75.52
FedTPG	68.51	77.05	72.53	FedTPG	73.76	77.05	75.37	FedTPG	74.34	77.05	75.67
TCP	75.51	78.10	76.78	TCP	76.09	79.15	77.59	TCP	79.88	78.31	79.09
PromptFolio	73.76	76.42	75.06	PromptFolio	79.59	75.57	77.53	PromptFolio	83.09	74.52	78.57
DP-FPL	69.38	<b>78.31</b>	73.58	DP-FPL	73.76	<b>79.57</b>	76.55	DP-FPL	75.80	75.78	75.79
GPR-NIAM <sub>P</sub>	82.50	75.57	78.89	GPR-NIAM <sub>P</sub>	84.83	78.31	81.44	GPR-NIAM <sub>P</sub>	85.13	77.47	81.12
GPR-NIAM <sub>T</sub>	87.11	77.83	<b>82.20</b>	GPR-NIAM <sub>T</sub>	<b>86.58</b>	78.94	<b>82.59</b>	GPR-NIAM <sub>T</sub>	<b>87.75</b>	<b>78.73</b>	<b>83.00</b>




Figure 4: Visualization of similarity heatmaps between text features and image patches before and after prompt refinement on Caltech101. In the similarity maps, red and yellow indicate higher similarity between image patches and the text features, while blue represents lower similarity.

## 5.5 Performance Comparison across Different Local Training Epoch Settings

To assess the robustness of the proposed GPR-NIAM under different local training efforts, we evaluate its performance with local epochs  $\in \{1, 5, 20\}$ , a data heterogeneity level of  $\beta = 0.5$ , and 10 participating clients. As shown in Table 4, both GPR-NIAM<sub>P</sub> and GPR-NIAM<sub>T</sub> consistently outperform baselines across all datasets and settings. Notably, GPR-NIAM<sub>T</sub> achieves the highest harmonic mean (HM) in most cases, indicating better generalization across base and novel classes. On the Flowers102 dataset, it maintains strong performance with increasing local epochs, reflecting robustness against overfitting. Overall, both variants remain stable and effective despite changing local training dynamics.

## 5.6 Effectiveness of Global Prompt Refinement

This section aims to evaluate the effectiveness of the global prompt refinement method. Specifically, we visualize the similarity maps between text features and image patches before and after refinement, as shown in Figure 4. Intuitively, the similarity responses between visual targets and textual semantics

are often scattered and misaligned before refinement. For example, in the “Dollar Bill” and “Bonsai” samples, the high similarity responses in the second row primarily concentrate on background regions near the image boundaries. In contrast, the responses become significantly more concentrated on semantically relevant regions, such as the face on the bill or the foliage of the bonsai after refinement. Notably, this advantage is also evident in novel classes. For instance, the refined prompts enable the model to better localize foreground objects while suppressing irrelevant background activations in “Kangaroo” and “Nautilus” samples. Moreover, it becomes tightly aligned with the spiral structure of the shell in the “Nautilus”, indicating enhanced semantic grounding.




Figure 5: Comparison between Random and Local Prompt-Tuning for Initialization.




Figure 6: Failure case analysis. Most GPR-NIAM errors overlap with zero-shot CLIP.

## 5.7 Evaluation of Local Prompt Tuning as Initialization

This section evaluates the performance of random initialization and local prompt tuning (LPT) on Flowers102, UCF101 and DTD datasets. As shown in Figure 5, LPT yields superior performance compared to random initialization on all cases. This can be attributed to the local training that can be more aligned with the target data, allowing the model to better capture domain-specific features. In particular, LPT outperforms the random initialization by 6% on the DTD dataset for base classes, which highlights the method’s ability to adapt more efficiently to the nuances of the texture.

## 5.8 Failure Case Analysis

This section reports the percentage of cases where the zero-shot CLIP model also failed in instances where the GPR-NIAM model made incorrect predictions (Figure 6). Clearly, there is a high overlap in the failure cases of both models. This reflects that the imprecision of zero-shot generalized knowledge hinders the model’s generalization, which also provides insights for future work to address the challenges posed by the lack of domain-specific adaptation in generalization models.

## 6 Conclusions and Future Work

In this paper, we balance the model capability to learn task-specific knowledge for downstream tasks with its capacity to retain generalized knowledge. Specifically, we propose GPR-NIAM for one-shot FPL, which introduces an attention masking mechanism to limit the interaction between the learnable prompt tokens and the original text tokens. Furthermore, it learns a unified prompt through global prompt refinement, thereby mitigating the data heterogeneity issue. Experimental results demonstrate that GPR-NIAM significantly improves the performance of global prompts across diverse visual tasks.

Despite the promising results of GPR-NIAM, there remains room for improvement. The communication overhead from prototype transmission may become a bottleneck in large-scale federated systems with limited bandwidth, which we plan to address in future work. And both the multi-label [68, 69], multi-modal [70, 71, 72, 73] problems and real-world medical imaging tasks [74, 75, 76] become considerably more challenging in federated settings. In addition, several issues discussed in Section 5.8 remain open and deserve further investigation.

## Acknowledgments

This work is supported in part by the Key Research and Development Program of Shandong Province (Grant No. 2024TSGC0667); the Ministry of Education, Singapore, under its Academic Research Fund Tier 1 (RG101/24); the RIE2025 Industry Alignment Fund – Industry Collaboration Projects (IAF-ICP) (Award I2301E0026), administered by A\*STAR, as well as supported by Alibaba Group and NTU Singapore through Alibaba-NTU Global e-Sustainability CorpLab (ANGEL).

## References

- [1] Zhuang Qi, Weihao He, Xiangxu Meng, and Lei Meng. Attentive modeling and distillation for out-of-distribution generalization of federated learning. In *2024 IEEE International Conference on Multimedia and Expo (ICME)*, pages 1–6. IEEE, 2024.
- [2] Jinyu Cai, Yunhe Zhan, Zhoumin Lu, Wenzhong Guo, and See-Kiong Ng. Towards effective federated graph anomaly detection via self-boosted knowledge distillation. In *Proceedings of the 32nd ACM International Conference on Multimedia*, pages 5537–5546, 2024.
- [3] Xianjie Guo, Kui Yu, Hao Wang, Lizhen Cui, Han Yu, and Xiaoxiao Li. Sample quality heterogeneity-aware federated causal discovery through adaptive variable space selection. In *Proceedings of the Thirty-Third International Joint Conference on Artificial Intelligence*, pages 4071–4079, 2024.
- [4] Tian Li, Anit Kumar Sahu, Ameet Talwalkar, and Virginia Smith. Federated learning: Challenges, methods, and future directions. *IEEE signal processing magazine*, 37(3):50–60, 2020.
- [5] Lele Fu, Bowen Deng, Sheng Huang, Tianchi Liao, Shirui Pan, and Chuan Chen. Less is more: Federated graph learning with alleviating topology heterogeneity from a causal perspective. In *Forty-second International Conference on Machine Learning*, 2025.
- [6] Jinyu Cai, Yunhe Zhang, Jicong Fan, and See-Kiong Ng. Lg-fgad: An effective federated graph anomaly detection framework. In *Proceedings of the International Joint Conference on Artificial Intelligence*, pages 3760–3769, 2024.
- [7] Xinghao Wu, Jianwei Niu, Xuefeng Liu, Tao Ren, Zhangmin Huang, and Zhetao Li. pfedgf: Enabling personalized federated learning via gradient fusion. In *2022 IEEE International Parallel and Distributed Processing Symposium (IPDPS)*, pages 639–649, 2022.
- [8] Liping Yi, Han Yu, Chao Ren, Gang Wang, Xiaoxiao Li, et al. Federated model heterogeneous matryoshka representation learning. *Advances in Neural Information Processing Systems*, 37:66431–66454, 2024.
- [9] Liping Yi, Han Yu, Gang Wang, and Xiaoguang Liu. pfedes: Model heterogeneous personalized federated learning with feature extractor sharing. *arXiv preprint arXiv:2311.06879*, 2023.
- [10] Alec Radford, Jong Wook Kim, Chris Hallacy, et al. Learning transferable visual models from natural language supervision. In *International conference on machine learning*, pages 8748–8763. PMLR, 2021.
- [11] Zhuang Qi, Sijin Zhou, Lei Meng, Han Hu, Han Yu, and Xiangxu Meng. Federated deconfounding and debiasing learning for out-of-distribution generalization. *arXiv preprint arXiv:2505.04979*, 2025.
- [12] Haodong Zhao, Wei Du, Fangqi Li, Peixuan Li, and Gongshen Liu. Fedprompt: Communication-efficient and privacy-preserving prompt tuning in federated learning. In *ICASSP 2023-2023 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 1–5. IEEE, 2023.
- [13] Chao Ren, Han Yu, et al. Advances and open challenges in federated foundation models. *IEEE Communications Surveys and Tutorials*, 2025.
- [14] Tao Fan, Hanlin Gu, et al. Ten challenging problems in federated foundation models. *IEEE Transactions on Knowledge and Data Engineering*, 2025.
- [15] Fu-En Yang, Chien-Yi Wang, and Yu-Chiang Frank Wang. Efficient model personalization in federated learning via client-specific prompt generation. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 19159–19168, 2023.
- [16] Xianjie Guo, Kui Yu, Lizhen Cui, Han Yu, and Xiaoxiao Li. Federated causally invariant feature learning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 16978–16986, 2025.
- [17] Bikang Pan et al. Federated learning from vision-language foundation models: Theoretical analysis and method. *Advances in Neural Information Processing Systems*, 37:30590–30623, 2024.
- [18] Wenke Huang, Mang Ye, Zekun Shi, Guancheng Wan, He Li, Bo Du, and Qiang Yang. Federated learning for generalization, robustness, fairness: A survey and benchmark. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2024.
- [19] Jiahua Dong, Lixu Wang, Zhen Fang, et al. Federated class-incremental learning. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2022.

- [20] Lele Fu, Sheng Huang, Yanyi Lai, Chuanfu Zhang, Hong-Ning Dai, Zibin Zheng, and Chuan Chen. Federated domain-independent prototype learning with alignments of representation and parameter spaces for feature shift. *IEEE Transactions on Mobile Computing*, 2025.
- [21] Jiayuan Zhang, Xuefeng Liu, Yukang Zhang, Guogang Zhu, Jianwei Niu, and Shaojie Tang. Enabling collaborative test-time adaptation in dynamic environment via federated learning. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 4191–4202, 2024.
- [22] Linh Tran, Wei Sun, Stacy Patterson, and Ana Milanova. Privacy-preserving personalized federated prompt learning for multimodal large language models. In *The International Conference on Learning Representations*, pages 1–19, 2025.
- [23] Tianyu Cui, Hongxia Li, Jingya Wang, and Ye Shi. Harmonizing generalization and personalization in federated prompt learning. *arXiv preprint arXiv:2405.09771*, 2024.
- [24] Hongxia Li, Wei Huang, Jingya Wang, and Ye Shi. Global and local prompts cooperation via optimal transport for federated learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 12151–12161, 2024.
- [25] Jie Zhang, Chen Chen, Bo Li, Lingjuan Lyu, Shuang Wu, Shouhong Ding, Chunhua Shen, and Chao Wu. Dense: Data-free one-shot federated learning. *Advances in Neural Information Processing Systems*, 35:21414–21428, 2022.
- [26] Hui Zeng, Minrui Xu, Tongqing Zhou, Xinyi Wu, Jiawen Kang, Zhiping Cai, and Dusit Niyato. One-shot-but-not-degraded federated learning. In *Proceedings of the 32nd ACM International Conference on Multimedia*, pages 11070–11079, 2024.
- [27] Yiqun Diao, Qinbin Li, and Bingsheng He. Towards addressing label skews in one-shot federated learning. In *The Eleventh International Conference on Learning Representations*, 2023.
- [28] Mingzhao Yang, Shangchao Su, Bin Li, and Xiangyang Xue. Exploring one-shot semi-supervised federated learning with pre-trained diffusion models. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 16325–16333, 2024.
- [29] Xinghao Wu, Jianwei Niu, Xuefeng Liu, Guogang Zhu, Shaojie Tang, Wanyu Lin, and Jiannong Cao. The diversity bonus: Learning from dissimilar clients in personalized federated learning. *IEEE Transactions on Neural Networks and Learning Systems*, pages 1–15, 2025.
- [30] Dezhong Yao, Yuexin Shi, Tongtong Liu, and Zhiqiang Xu. Fedmho: Heterogeneous one-shot federated learning towards resource-constrained edge devices. *arXiv preprint arXiv:2502.08518*, 2025.
- [31] Yanbiao Ma, Wei Dai, Wenke Huang, and Jiayi Chen. Geometric knowledge-guided localized global distribution alignment for federated learning. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pages 20958–20968, 2025.
- [32] Ming Hu, Yue Cao, Anran Li, Zhiming Li, Chengwei Liu, Tianlin Li, Mingsong Chen, and Yang Liu. Fedmut: Generalized federated learning via stochastic mutation. In *Proceedings of the AAAI conference on artificial intelligence*, volume 38, pages 12528–12537, 2024.
- [33] Shangchao Su, Mingzhao Yang, Bin Li, and Xiangyang Xue. Federated adaptive prompt tuning for multi-domain collaborative learning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 15117–15125, 2024.
- [34] Yongju Jia, Jiarui Ma, Xiangxian Li, Baiqiao Zhang, Xianhui Cao, Juan Liu, and Yulong Bian. Llm-empowered dynamic prompt routing for vision-language models tuning under long-tailed distributions. *arXiv preprint arXiv:2508.15688*, 2025.
- [35] Tao Guo, Song Guo, and Junxiao Wang. Pfedprompt: Learning personalized prompt for vision-language models in federated learning. In *Proceedings of the ACM Web Conference 2023*, pages 1364–1374, 2023.
- [36] Wenlong Deng, Christos Thrampoulidis, and Xiaoxiao Li. Unlocking the potential of prompt-tuning in bridging generalized and personalized federated learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 6087–6097, 2024.
- [37] Yichen Li, Xiuying Wang, Wenchoao Xu, Haozhao Wang, Yining Qi, Jiahua Dong, and Ruixuan Li. Feature distillation is the better choice for model-heterogeneous federated learning. *arXiv preprint arXiv:2507.10348*, 2025.

- [38] Mingzhao Yang, Shangchao Su, Bin Li, and Xiangyang Xue. Feddeo: Description-enhanced one-shot federated learning with diffusion models. In *Proceedings of the 32nd ACM International Conference on Multimedia*, pages 6666–6675, 2024.
- [39] Youssef Allouah, Akash Dhasade, Rachid Guerraoui, Nirupam Gupta, Anne-Marie Kermarrec, Rafael Pinot, Rafael Pires, and Rishi Sharma. Revisiting ensembling in one-shot federated learning. *arXiv preprint arXiv:2411.07182*, 2024.
- [40] Yi-Han Chiang, Koudai Terai, Tsung-Wei Chiang, Hai Lin, Yusheng Ji, and John CS Lui. Optimal transport-based one-shot federated learning for artificial intelligence of things. *IEEE Internet of Things Journal*, 11(2):2166–2180, 2023.
- [41] Clare Elizabeth Heinbaugh, Emilio Luz-Ricca, and Huajie Shao. Data-free one-shot federated learning under very high statistical heterogeneity. In *The Eleventh International Conference on Learning Representations*, 2023.
- [42] Naibo Wang, Yuchen Deng, Wenjie Feng, Shichen Fan, Jianwei Yin, and See-Kiong Ng. One-shot sequential federated learning for non-iid data by enhancing local model diversity. In *Proceedings of the 32nd ACM International Conference on Multimedia*, pages 5201–5210, 2024.
- [43] Ming Hu, Peiheng Zhou, Zhihao Yue, Zhiwei Ling, Yihao Huang, Anran Li, Yang Liu, Xiang Lian, and Mingsong Chen. Fedcross: Towards accurate federated learning via multi-model cross-aggregation. In *2024 IEEE 40th International Conference on Data Engineering (ICDE)*, pages 2137–2150. IEEE, 2024.
- [44] Zhuang Qi, Runhui Zhang, Lei Meng, Wei Wu, Yachong Zhang, and Xiangxu Meng. Global intervention and distillation for federated out-of-distribution generalization. *arXiv preprint arXiv:2504.00850*, 2025.
- [45] Yichen Li, Haozhao Wang, Yining Qi, Wei Liu, and Ruixuan Li. Re-fed+: A better replay strategy for federated incremental learning. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2025.
- [46] Jiayuan Zhang, Xuefeng Liu, Jianwei Niu, et al. Causality inspired federated learning for ood generalization. In *Forty-second International Conference on Machine Learning*.
- [47] Zhuang Qi, Lei Meng, Zitan Chen, Han Hu, Hui Lin, and Xiangxu Meng. Cross-silo prototypical calibration for federated learning with non-iid data. In *Proceedings of the 31st ACM International Conference on Multimedia*, pages 3099–3107, 2023.
- [48] Lei Meng, Zhuang Qi, Lei Wu, Xiaoyu Du, Zhaochuan Li, Lizhen Cui, and Xiangxu Meng. Improving global generalization and local personalization for federated learning. *IEEE Transactions on Neural Networks and Learning Systems*, 36(1):76–87, 2025.
- [49] Zhuang Qi, Lei Meng, Weihao He, et al. Cross-training with multi-view knowledge fusion for heterogenous federated learning. *arXiv preprint arXiv:2405.20046*, 2024.
- [50] Xiangyu Meng, Xiangxian Li, Lin Zhao, Yajuan Shen, Hui Sun, Xiaoming Cong, Xianhui Cao, Yunfeng Bi, Juan Liu, and Yulong Bian. Robust representation of domain incomplete tabular data via cross-modal debiasing and prototype-fused reconstruction. *Journal of King Saud University Computer and Information Sciences*, 37(7):209, 2025.
- [51] Zhuang Qi, Lei Meng, Zhaohcuan Li, Han Hu, and Xiangxu Meng. Cross-silo feature space alignment for federated learning on clients with imbalanced data. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 1–9, 2025.
- [52] Chen Qiu, Xingyu Li, Chaithanya Kumar Mumadi, Madan Ravi Ganesh, Zhenzhen Li, Lu Peng, and Wan-Yi Lin. Federated text-driven prompt generation for vision-language models. In *The Twelfth International Conference on Learning Representations*, 2024.
- [53] Alex Krizhevsky, Geoffrey Hinton, et al. Learning multiple layers of features from tiny images. 2009.
- [54] Omkar M Parkhi, Andrea Vedaldi, Andrew Zisserman, and CV Jawahar. Cats and dogs. In *2012 IEEE conference on computer vision and pattern recognition*, pages 3498–3505. IEEE, 2012.
- [55] Li Fei-Fei, Rob Fergus, and Pietro Perona. Learning generative visual models from few training examples: An incremental bayesian approach tested on 101 object categories. In *2004 conference on computer vision and pattern recognition workshop*, pages 178–178. IEEE, 2004.
- [56] Mircea Cimpoi, Subhransu Maji, Iasonas Kokkinos, Sammy Mohamed, and Andrea Vedaldi. Describing textures in the wild. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 3606–3613, 2014.

- [57] Subhransu Maji, Esa Rahtu, Juho Kannala, Matthew Blaschko, and Andrea Vedaldi. Fine-grained visual classification of aircraft. *arXiv preprint arXiv:1306.5151*, 2013.
- [58] Maria-Elena Nilsback and Andrew Zisserman. Automated flower classification over a large number of classes. In *2008 Sixth Indian conference on computer vision, graphics & image processing*, pages 722–729. IEEE, 2008.
- [59] Jonathan Krause, Michael Stark, Jia Deng, and Li Fei-Fei. 3d object representations for fine-grained categorization. In *Proceedings of the IEEE international conference on computer vision workshops*, pages 554–561, 2013.
- [60] Khurram Soomro, Amir Roshan Zamir, and Mubarak Shah. Ucf101: A dataset of 101 human actions classes from videos in the wild. *arXiv preprint arXiv:1212.0402*, 2012.
- [61] Antonio Torralba and Alexei A Efros. Unbiased look at dataset bias. In *CVPR 2011*, pages 1521–1528. IEEE, 2011.
- [62] Xingchao Peng, Qinxun Bai, Xide Xia, Zijun Huang, Kate Saenko, and Bo Wang. Moment matching for multi-source domain adaptation. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 1406–1415, 2019.
- [63] Haoyuan Liang, Xinyu Zhang, Shilei Cao, Guowen Li, and Juepeng Zheng. Tta-feddg: Leveraging test-time adaptation to address federated domain generalization. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 18658–18666, 2025.
- [64] Wenhao Li, Qiangchang Wang, Peng Zhao, and Yilong Yin. Knn transformer with pyramid prompts for few-shot learning. In *Proceedings of the 32nd ACM International Conference on Multimedia*, pages 1082–1091, 2024.
- [65] Tao Guo, Song Guo, Junxiao Wang, Xueyang Tang, and Wencho Xu. Promptfl: Let federated participants cooperatively learn prompts instead of models–federated learning in age of foundation model. *IEEE Transactions on Mobile Computing*, 23(5):5179–5194, 2023.
- [66] Dongjun Lee, Seokwon Song, Jihee Suh, Joonmyeong Choi, Sanghyeon Lee, and Hyunwoo J Kim. Read-only prompt optimization for vision-language few-shot learning. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 1401–1411, 2023.
- [67] Hantao Yao, Rui Zhang, and Changsheng Xu. Tcp: Textual-based class-aware prompt tuning for visual-language model. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 23438–23448, 2024.
- [68] Zhiqiang Kou, Si Qin, Hailin Wang, Mingkun Xie, et al. Label distribution learning with biased annotations by learning multi-label representation, 2025.
- [69] Zhiqiang Kou, Jing Wang, Jiawei Tang, Yuheng Jia, Boyu Shi, and Xin Geng. Exploiting multi-label correlation in label distribution learning. In Kate Larson, editor, *Proceedings of the Thirty-Third International Joint Conference on Artificial Intelligence, IJCAI-24*, pages 4326–4334. International Joint Conferences on Artificial Intelligence Organization, 8 2024. Main Track.
- [70] Lei Meng, Xiangxian Li, Xiaoshuo Yan, Haokai Ma, Zhuang Qi, Wei Wu, and Xiangxu Meng. Causal inference over visual-semantic-aligned graph for image classification. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 19449–19457, 2025.
- [71] Zheng Wang, Zhenwei Gao, Yang Yang, et al. Geometric matching for cross-modal retrieval. *IEEE Transactions on Neural Networks and Learning Systems*, 36(3):5509–5521, 2025.
- [72] Yuqing Wang, Zhuang Qi, Xiangxian Li, Jinxing Liu, Xiangxu Meng, and Lei Meng. Multi-channel attentive weighting of visual frames for multimodal video classification. In *2023 International Joint Conference on Neural Networks (IJCNN)*, pages 1–8. IEEE, 2023.
- [73] Harry Cheng, Yangyang Guo, Tianyi Wang, et al. Voice-face homogeneity tells deepfake. *ACM Transactions on Multimedia Computing, Communications and Applications*, 20(3):1–22, 2023.
- [74] Wei Feng, Lie Ju, Lin Wang, Kaimin Song, Xin Zhao, and Zongyuan Ge. Unsupervised domain adaptation for medical image segmentation by selective entropy constraints and adaptive semantic alignment. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37, pages 623–631, 2023.
- [75] Wei Feng, Sijin Zhou, et al. Neighbor-guided unbiased framework for generalized category discovery in medical image classification. *IEEE Journal of Biomedical and Health Informatics*, 2025.
- [76] Chengrun Dang, Zhuang Qi, Tao Xu, Mingkai Gu, et al. Deep learning-powered whole slide image analysis in cancer pathology. *Laboratory Investigation*, page 104186, 2025.

## A Appendix

### A.1 Datasets

To conduct our evaluation, we selected 10 diverse visual classification datasets as benchmarks. Table 5 provides a detailed overview, including the number of classes, the size of training and testing sets, the number of domains, and the corresponding federated settings.

For datasets with multiple domains, we adopt the commonly used Office-Home benchmark, which includes four domains: Art, Clipart, Product, and RealWorld. In addition, we utilize DomainNet, a large-scale dataset spanning six domains: Clipart, Infograph, Painting, Quickdraw, Real, and Sketch. We focus on training with a selected subset of 20 classes from each dataset. Figure 7 illustrates representative raw instances from the two multi-domain datasets.

Table 5: Statistics and federated settings of the datasets used in our experiments.

Dataset	Classes	Train	Test		Domains	Federated Settings	
			Base	Novel		Clients	Heterogeneity
CIFAR10	10	25,000	5,000	5,000	1	5/10/20	0.1/0.5
OxfordPets	37	1,785	1,782	1,887	1	5/10/20	0.1/0.5
Caltech101	100	3,310	1,416	1,178	1	5/10/20	0.1/0.5
DTD	47	1,440	864	828	1	5/10/20	0.1/0.5
FGVCAircraft	100	3,333	1,667	1,666	1	5/10/20	0.1/0.5
Flowers102	102	2,807	343	475	1	5/10/20	0.1/0.5
StanfordCars	196	4,052	4,002	3,998	1	5/10/20	0.1/0.5
UCF101	101	3,926	1,895	1,888	1	5/10/20	0.1/0.5
DomainNet	20		30,777		6	9/15	–
Office-Home	65		15,588		4	15/25	–




Figure 7: Illustrative examples of class instances across domains: “Ambulance” from DomainNet and “Bed” from Office-Home.

### A.2 Experimental Results

#### A.2.1 Base-to-Base/Novel Generalization under High-Level Data Heterogeneity

To more comprehensively evaluate the effectiveness of the methods, we test all approaches on eight datasets under a higher level of heterogeneity ( $\beta = 0.1$ ), focusing on their performance in the base-to-base/novel generalization setting. As presented in Table 6, our proposed method GPR-NIAM obtains the highest average

Table 6: Performance comparison between GPR-NIAM and baselines on 8 datasets with  $\beta = 0.1$  in the base-to-base/novel generalization setting. The best performance is indicated in bold.

(a) Average over 8 datasets				(b) CIFAR10				(c) OxfordPets			
Method	Base	Novel	HM	Method	Base	Novel	HM	Method	Base	Novel	HM
CLIP	65.79	74.32	69.70	CLIP	95.30	95.50	95.39	CLIP	81.53	95.33	87.89
PromptFL	69.15	72.50	70.71	PromptFL	94.72	95.68	95.19	PromptFL	89.33	95.91	92.51
RPO	67.81	73.66	70.57	RPO	95.10	94.76	94.92	RPO	86.19	95.54	90.63
FedTPG	67.06	71.71	69.20	FedTPG	94.30	96.64	95.45	FedTPG	86.58	94.70	90.46
TCP	70.30	74.19	72.13	TCP	95.30	95.66	95.47	TCP	90.01	<b>96.23</b>	93.02
PromptFolio	69.49	73.44	71.33	PromptFolio	95.20	96.74	95.96	PromptFolio	90.90	95.17	92.99
DP-FPL	71.90	71.75	69.88	DP-FPL	94.24	95.42	94.82	DP-FPL	91.30	95.12	<b>93.17</b>
FedDISC	52.27	54.40	50.47	FedDISC	92.47	91.37	91.91	FedDISC	81.48	66.74	73.38
FedELMY	68.05	47.29	53.41	FedELMY	89.84	28.90	43.73	FedELMY	86.70	82.77	84.69
GPR-NIAM <sub>P</sub>	<b>73.87</b>	72.84	73.18	GPR-NIAM <sub>P</sub>	95.88	<b>96.84</b>	<b>96.35</b>	GPR-NIAM <sub>P</sub>	90.57	95.23	92.84
GPR-NIAM <sub>T</sub>	73.62	<b>74.48</b>	<b>73.97</b>	GPR-NIAM <sub>T</sub>	<b>96.03</b>	96.24	96.13	GPR-NIAM <sub>T</sub>	<b>91.86</b>	94.38	93.10
(d) Caltech101				(e) DTD				(f) FGVCAircraft			
Method	Base	Novel	HM	Method	Base	Novel	HM	Method	Base	Novel	HM
CLIP	82.27	94.05	87.77	CLIP	53.35	59.05	56.06	CLIP	26.51	31.21	28.67
PromptFL	87.07	93.03	89.95	PromptFL	56.36	52.53	54.38	PromptFL	26.03	30.61	28.13
RPO	84.95	93.80	89.16	RPO	53.12	<b>60.02</b>	56.36	RPO	26.51	30.37	28.31
FedTPG	85.38	93.37	89.20	FedTPG	53.35	50.00	51.62	FedTPG	21.17	26.35	23.48
TCP	88.41	<b>94.16</b>	91.19	TCP	58.10	57.60	57.85	TCP	28.37	32.23	30.18
PromptFolio	82.41	92.61	87.21	PromptFolio	53.58	55.55	54.55	PromptFolio	28.13	<b>34.15</b>	<b>30.85</b>
DP-FPL	87.57	93.54	90.46	DP-FPL	51.85	45.65	48.55	DP-FPL	26.63	29.35	27.92
FedDISC	22.37	81.39	35.09	FedDISC	65.23	41.23	49.87	FedDISC	18.49	17.61	18.03
FedELMY	85.59	90.40	87.93	FedELMY	40.27	11.95	18.43	FedELMY	23.45	5.46	8.86
GPR-NIAM <sub>P</sub>	<b>91.24</b>	93.88	<b>92.54</b>	GPR-NIAM <sub>P</sub>	<b>67.59</b>	52.89	59.34	GPR-NIAM <sub>P</sub>	<b>29.69</b>	30.81	30.23
GPR-NIAM <sub>T</sub>	90.60	94.05	92.30	GPR-NIAM <sub>T</sub>	65.62	57.79	<b>61.45</b>	GPR-NIAM <sub>T</sub>	29.39	29.96	29.67
(g) Flowers102				(h) StanfordCars				(i) UCF101			
Method	Base	Novel	HM	Method	Base	Novel	HM	Method	Base	Novel	HM
CLIP	62.39	78.10	69.36	CLIP	56.87	69.40	62.51	CLIP	68.12	71.92	69.97
PromptFL	72.88	75.78	74.30	PromptFL	57.31	67.58	62.02	PromptFL	69.55	68.90	69.22
RPO	72.01	76.63	74.24	RPO	57.64	66.73	61.85	RPO	67.01	71.45	69.16
FedTPG	70.84	77.26	73.91	FedTPG	56.89	<b>69.55</b>	62.59	FedTPG	68.02	65.83	66.91
TCP	73.61	77.94	75.71	TCP	57.79	69.28	63.01	TCP	70.88	70.43	70.65
PromptFolio	75.51	76.42	75.96	PromptFolio	58.97	69.05	63.61	PromptFolio	71.29	67.84	69.52
DP-FPL	68.97	71.43	70.08	DP-FPL	53.99	63.65	58.43	DP-FPL	73.61	69.80	71.66
FedDISC	25.44	32.47	28.52	FedDISC	42.59	51.23	46.51	FedDISC	70.12	53.23	60.51
FedELMY	<b>81.63</b>	57.05	67.16	FedELMY	60.26	49.94	54.62	FedELMY	<b>76.72</b>	51.90	61.92
GPR-NIAM <sub>P</sub>	79.30	77.26	78.26	GPR-NIAM <sub>P</sub>	<b>60.81</b>	66.98	63.75	GPR-NIAM <sub>P</sub>	75.88	68.85	72.19
GPR-NIAM <sub>T</sub>	80.17	<b>79.15</b>	<b>79.65</b>	GPR-NIAM <sub>T</sub>	60.14	68.65	<b>64.12</b>	GPR-NIAM <sub>T</sub>	75.19	<b>75.63</b>	<b>75.41</b>

performance on the eight datasets. Specifically, GPR-NIAM<sub>P</sub> and GPR-NIAM<sub>T</sub> achieve the top two performances on the base classes, ranking first and second, respectively. In terms of novel class generalization, GPR-NIAM<sub>T</sub> outperforms all baselines, demonstrating its superior capability to adapt to unseen categories. Furthermore, GPR-NIAM<sub>T</sub> also achieves the highest overall performance in terms of harmonic mean (HM), indicating a strong balance between knowledge retention and generalization. These results highlight the robustness and adaptability of the proposed method under highly heterogeneous federated settings.

### A.2.2 Leave-One-Domain-Out Generalization with Increased Client Diversity

To validate the effectiveness of the proposed GPR-NIAM under more challenging scenarios, we conduct experiments in a leave-one-domain-out generalization setting [] using the Office-Home and DomainNet datasets. In this setting, the data from each domain is randomly partitioned among 5 clients, and for each evaluation, one domain is held out entirely for testing while the remaining domains are used for federated training. As shown in Table 7, GPR-NIAM<sub>P</sub> and GPR-NIAM<sub>T</sub> consistently outperform existing baselines across both datasets. For the Office-Home, GPR-NIAM<sub>T</sub> achieves the highest average accuracy, with GPR-NIAM<sub>P</sub> closely following. Notably, GPR-NIAM<sub>T</sub> yields the best performance on the “RealWorld” and “Product” domains, while GPR-NIAM<sub>P</sub> leads on “Art” and “Clipart”. For the DomainNet, both variants of GPR-NIAM also demonstrate strong cross-domain

Table 7: Performance comparison between GPR-NIAM and baselines on the Office-Home and DomainNet datasets under the leave-one-domain-out generalization setting, where the data from each domain is randomly partitioned among 5 clients.

Method	Office-Home					DomainNet						
	Art	Clipart	Product	RealWorld	Average	Clipart	Infograph	Painting	Quickdraw	Real	Sketch	Average
CLIP	74.57	62.33	80.17	79.38	74.08	92.87	77.02	85.35	43.06	98.61	94.22	81.85
PromptFL	73.71	64.49	82.81	80.39	75.35	94.93	80.09	<b>87.94</b>	52.68	99.00	95.63	85.04
RPO	73.83	62.13	79.95	79.36	73.81	93.66	77.31	84.65	44.46	98.83	93.87	82.13
FedTPG	74.90	65.54	82.85	80.37	75.91	94.35	79.76	87.40	48.05	98.92	<b>95.89</b>	84.06
TCP	75.77	65.10	83.10	80.99	76.24	93.92	79.72	87.67	46.82	98.96	95.00	83.68
PromptFolio	75.81	65.36	83.53	80.39	76.27	<b>95.51</b>	<b>81.00</b>	89.33	48.82	98.99	95.20	84.80
DP-FPL	74.82	63.55	82.85	78.79	75.00	94.77	80.20	89.00	48.35	98.78	94.82	84.32
FedDISC	71.19	62.48	81.31	78.35	73.48	94.23	80.25	83.85	44.21	98.39	92.49	82.23
FedELMY	75.45	64.29	82.36	77.58	74.92	94.15	75.48	86.89	47.39	98.72	95.13	82.96
GPR-NIAM <sub>P</sub>	<b>78.16</b>	<b>72.76</b>	89.09	83.68	80.92	94.45	80.60	87.50	<b>58.68</b>	<b>99.10</b>	94.91	85.87
GPR-NIAM <sub>T</sub>	77.66	71.95	<b>89.47</b>	<b>85.93</b>	<b>81.25</b>	95.19	80.49	87.23	57.96	99.07	95.37	<b>85.88</b>

generalization, achieving top results in multiple domains. These results highlight the robustness of our method in handling both domain shift and non-IID data across clients. In summary, the proposed GPR-NIAM exhibits superior generalization performance under leave-one-domain-out conditions, demonstrating its effectiveness in federated scenarios with high domain and client-level variability.

### A.2.3 Performance Comparison across Different Local Training Epoch Settings

Table 8: Performance comparison between GPR-NIAM and baselines on StanfordCars and UCF101 with local training epochs  $\in \{1, 5, 20\}$  in the base-to-base/novel generalization setting.

(g) StanfordCars (Epochs=1)			(h) StanfordCars (Epochs=5)			(i) StanfordCars (Epochs=20)					
Method	Base	Novel	HM	Method	Base	Novel	HM	Method	Base	Novel	HM
CLIP	56.87	69.40	62.51	CLIP	56.87	69.40	62.51	CLIP	56.87	69.40	62.51
PromptFL	58.87	68.40	63.27	PromptFL	60.39	67.35	63.68	PromptFL	59.74	67.83	63.52
RPO	56.39	69.03	62.07	RPO	57.12	69.70	62.79	RPO	57.42	<b>70.06</b>	63.11
FedTPG	57.09	69.73	62.78	FedTPG	58.14	69.65	63.38	FedTPG	58.69	69.25	63.54
TCP	58.04	69.40	63.22	TCP	58.72	69.91	63.82	TCP	59.29	68.68	63.64
PromptFolio	58.69	<b>69.83</b>	63.78	PromptFolio	60.04	69.78	64.54	PromptFolio	61.46	69.10	65.06
DP-FPL	49.30	62.85	55.25	DP-FPL	59.02	68.33	63.33	DP-FPL	59.24	64.35	61.69
GPR-NIAM <sub>P</sub>	60.61	67.58	63.91	GPR-NIAM <sub>P</sub>	62.09	67.05	64.48	GPR-NIAM <sub>P</sub>	<b>64.46</b>	66.38	65.41
GPR-NIAM <sub>T</sub>	<b>61.54</b>	67.40	64.33	GPR-NIAM <sub>T</sub>	60.99	69.45	<b>64.94</b>	GPR-NIAM <sub>T</sub>	63.01	69.00	<b>65.87</b>
(j) UCF101 (Epochs=1)			(k) UCF101 (Epochs=5)			(l) UCF101 (Epochs=20)					
Method	Base	Novel	HM	Method	Base	Novel	HM	Method	Base	Novel	HM
CLIP	68.12	71.92	69.97	CLIP	68.12	71.92	69.97	CLIP	68.12	71.92	69.97
PromptFL	69.55	69.01	69.28	PromptFL	73.03	72.08	72.55	PromptFL	68.07	64.24	66.10
RPO	65.96	70.55	68.17	RPO	66.80	71.50	69.07	RPO	68.75	71.13	69.92
FedTPG	64.69	64.08	64.39	FedTPG	68.86	67.47	68.16	FedTPG	68.86	66.73	67.77
TCP	73.82	72.47	73.13	TCP	75.83	72.27	74.00	TCP	76.09	73.00	74.51
PromptFolio	73.72	69.80	71.71	PromptFolio	76.56	70.12	73.20	PromptFolio	78.15	68.37	72.94
DP-FPL	72.13	70.65	71.38	DP-FPL	74.08	71.76	72.91	DP-FPL	73.24	68.16	70.61
GPR-NIAM <sub>P</sub>	76.88	68.53	72.47	GPR-NIAM <sub>P</sub>	<b>79.10</b>	68.75	73.56	GPR-NIAM <sub>P</sub>	<b>78.89</b>	69.38	73.83
GPR-NIAM <sub>T</sub>	<b>77.78</b>	<b>74.62</b>	<b>76.17</b>	GPR-NIAM <sub>T</sub>	78.15	<b>75.79</b>	<b>76.95</b>	GPR-NIAM <sub>T</sub>	78.04	<b>76.58</b>	<b>77.31</b>

To evaluate the effectiveness of the proposed GPR-NIAM under different levels of local training effort, this section investigates its performance across varying local training epochs. Specifically, we consider three settings with local training epochs  $\in \{1, 5, 20\}$ , under a data heterogeneity level ( $\beta = 0.5$ ) and with 10 clients participating in training. As shown in Table 4 and Table 8, both GPR-NIAM<sub>P</sub> and GPR-NIAM<sub>T</sub> consistently outperform the baselines across all datasets and training epochs. Notably, GPR-NIAM<sub>T</sub> delivers the best overall performance in terms of harmonic mean (HM) in most settings, which demonstrates superior generalization across both base and novel classes. On Flowers102 and UCF101, GPR-NIAM<sub>T</sub> shows strong robustness even when the number of local epochs increases. This highlights its ability to prevent overfitting to local distributions. Across datasets, both variants of GPR-NIAM remain stable and effective even as the local training dynamics change, outperforming methods that are more sensitive to overfitting or class imbalance under prolonged local updates.

#### A.2.4 Sensitivity Analysis Regarding the Number of Prompt Tokens.

Table 9: Comparison of different methods under two settings ( $5 \times 512$  and  $20 \times 512$ ) on Caltech101, DTD, and UCF101 datasets. Bold numbers indicate the best performance.

	Caltech101			DTD			UCF101		
	Base	Novel	HM	Base	Novel	HM	Base	Novel	HM
<b>5x512</b>									
PromptFL	88.41	<b>93.97</b>	91.11	59.25	52.77	55.83	76.78	69.91	73.18
FedTPG	86.65	93.37	89.88	60.18	56.15	58.10	74.30	70.63	72.41
PromptFolio	89.47	93.46	91.42	<b>61.02</b>	<b>56.28</b>	57.62	74.40	<b>70.65</b>	72.48
GRP-NIAM_P	<b>89.90</b>	93.37	<b>91.60</b>	60.96	55.95	<b>59.12</b>	<b>77.98</b>	70.39	<b>73.99</b>
<b>20x512</b>									
PromptFL	88.91	<b>94.31</b>	91.53	<b>62.50</b>	52.53	57.08	76.83	69.12	72.77
FedTPG	88.84	92.86	90.81	58.10	54.83	56.41	76.30	73.41	74.41
PromptFolio	89.97	94.22	92.05	56.13	53.98	55.03	75.72	74.41	75.06
GRP-NIAM_P	<b>90.02</b>	94.28	<b>92.10</b>	61.43	<b>56.23</b>	<b>58.71</b>	<b>77.23</b>	<b>74.58</b>	<b>75.88</b>

We have conducted additional sensitivity experiments regarding the number of prompt tokens. Specifically, we adjusted the prompt length to  $5 \times 512$  and  $20 \times 512$ . The results show that GPR-NIAM remains robust across different prompt lengths, while consistently maintaining an advantage over all baselines in all case.

#### A.2.5 Evaluating Method Scalability with Varying Client Numbers

Table 10: Performance comparison between GPR-NIAM and baselines on DTD and Flowers102 datasets under varying number of clients  $K \in \{20, 30\}$  in the base-to-base/novel generalization.

(a) DTD (K=20)			(b) DTD (K=30)				
Method	Base	Novel	HM	Method	Base	Novel	HM
CLIP	53.35	59.05	56.06	CLIP	53.35	59.05	56.06
PromptFL	54.16	54.22	54.19	PromptFL	52.89	56.52	54.64
RPO	53.24	<b>60.02</b>	56.42	RPO	52.77	<b>59.66</b>	56.00
FedTPG	57.83	54.93	56.34	FedTPG	52.19	58.09	54.98
TCP	58.21	55.93	57.04	TCP	56.82	58.42	57.60
PromptFolio	58.18	52.91	55.41	PromptFolio	55.43	52.41	53.88
DP-FPL	56.13	50.72	53.29	DP-FPL	57.06	51.20	53.97
GRP-NIAM <sub>P</sub>	<b>63.31</b>	52.29	57.27	GRP-NIAM <sub>P</sub>	<b>60.64</b>	51.93	55.95
GRP-NIAM <sub>T</sub>	62.38	53.86	<b>57.81</b>	GRP-NIAM <sub>T</sub>	60.41	56.52	<b>58.40</b>
(d) Flowers102 (K=20)			(e) Flowers102 (K=30)				
Method	Base	Novel	HM	Method	Base	Novel	HM
CLIP	62.39	78.10	69.36	CLIP	62.39	78.10	69.36
PromptFL	74.34	74.10	74.22	PromptFL	72.01	74.73	73.34
RPO	71.72	76.84	74.19	RPO	71.13	76.42	73.68
FedTPG	74.63	<b>78.31</b>	76.43	FedTPG	74.92	78.31	76.58
TCP	75.21	77.68	76.43	TCP	73.46	78.73	76.01
PromptFolio	75.51	<b>78.31</b>	76.88	PromptFolio	73.76	78.31	75.97
DP-FPL	74.34	74.10	74.22	DP-FPL	69.09	75.57	72.19
GRP-NIAM <sub>P</sub>	82.79	77.05	79.82	GRP-NIAM <sub>P</sub>	81.63	78.31	79.93
GRP-NIAM <sub>T</sub>	<b>83.96</b>	77.05	<b>80.36</b>	GRP-NIAM <sub>T</sub>	<b>82.50</b>	<b>78.94</b>	<b>80.68</b>

This section aims to evaluate the scalability of the proposed GPR-NIAM with respect to the number of participating clients. We consider two settings with  $K \in \{20, 30\}$  under a heterogeneity level  $\beta = 0.5$  and fixed local training epochs of 5. As shown in Table 10 and Table 11, both GPR-NIAM<sub>P</sub> and GPR-NIAM<sub>T</sub> consistently outperform all baselines across datasets and client configurations. GPR-NIAM<sub>T</sub> achieves the best overall performance on most datasets and demonstrates strong generalization to novel classes. For example, on Flowers102 and UCF101, it achieves the highest harmonic mean (HM) under both client settings. Meanwhile, GPR-NIAM<sub>P</sub> maintains strong

Table 11: Performance comparison between GPR-NIAM and baselines on StanfordCars and UCF101 datasets under varying number of clients  $K \in \{20, 30\}$  in the base-to-base/novel generalization.

(g) StanfordCars (K=20)			(h) StanfordCars (K=30)				
Method	Base	Novel	HM	Method	Base	Novel	HM
CLIP	56.87	69.40	62.51	CLIP	56.87	69.40	62.51
PromptFL	59.82	67.03	63.22	PromptFL	54.42	66.33	59.79
RPO	57.22	<b>69.65</b>	62.83	RPO	57.24	<b>69.60</b>	62.82
FedTPG	59.24	66.73	62.76	FedTPG	59.42	69.08	63.88
TCP	58.77	67.78	62.95	TCP	59.32	68.21	63.45
PromptFolio	60.91	65.78	63.25	PromptFolio	59.59	65.25	62.29
DP-FPL	57.74	65.65	61.44	DP-FPL	58.37	67.75	62.71
GPR-NIAM <sub>P</sub>	60.04	68.23	63.87	GPR-NIAM <sub>P</sub>	60.09	67.40	63.54
GPR-NIAM <sub>T</sub>	<b>61.02</b>	69.40	<b>64.94</b>	GPR-NIAM <sub>T</sub>	<b>60.14</b>	69.38	<b>64.43</b>
(j) UCF101 (K=20)			(k) UCF101 (K=30)				
Method	Base	Novel	HM	Method	Base	Novel	HM
CLIP	68.12	71.92	69.97	CLIP	68.12	71.92	69.97
PromptFL	75.67	68.75	72.04	PromptFL	71.34	66.04	68.59
RPO	66.80	71.34	69.00	RPO	66.75	71.39	68.99
FedTPG	74.88	74.78	74.83	FedTPG	76.09	74.62	75.35
TCP	74.08	73.83	73.96	TCP	73.50	75.58	74.53
PromptFolio	74.45	71.76	73.08	PromptFolio	77.09	74.04	75.54
DP-FPL	75.51	69.75	72.52	DP-FPL	72.71	68.96	70.79
GPR-NIAM <sub>P</sub>	76.41	69.01	72.52	GPR-NIAM <sub>P</sub>	76.62	72.50	74.50
GPR-NIAM <sub>T</sub>	<b>77.04</b>	<b>76.21</b>	<b>76.62</b>	GPR-NIAM <sub>T</sub>	<b>77.62</b>	<b>76.05</b>	<b>76.83</b>

performance on base classes. These results show that GPR-NIAM maintains high performance and stability even as the number of clients increases, demonstrating its scalability and effectiveness in more challenging cases.

#### A.2.6 Evaluation of Module Contributions

Table 12: Ablation study on the effectiveness of NIAM and CSCR modules under the base-to-base/novel generalization. Results are reported on four datasets with PromptFL as the baseline.

(a) DTD			(b) Flowers102				
Method	Base	Novel	HM	Method	Base	Novel	HM
PromptFL	55.55	56.76	56.15	PromptFL	78.42	76.00	77.19
+NIAM	55.78	54.73	55.25	+NIAM	76.13	77.78	76.94
+NIAM+CSCR	68.51	52.65	59.54	+NIAM+CSCR	84.83	78.31	81.44
(c) StanfordCars			(d) UCF101				
Method	Base	Novel	HM	Method	Base	Novel	HM
PromptFL	60.39	67.35	63.68	PromptFL	73.03	72.08	72.55
+NIAM	55.17	68.13	60.96	+NIAM	67.96	73.02	70.39
+NIAM+CSCR	62.09	67.05	64.48	+NIAM+CSCR	79.10	68.75	73.56

We evaluate the contributions of the NIAM and CSCR modules by gradually adding them to the PromptFL baseline. Experiments are conducted on four datasets using 10 clients, with heterogeneity level  $\beta = 0.5$  and 5 local training epochs. As shown in Table 12, NIAM improves performance on novel classes by preserving generalized knowledge. CSCR enhances base class performance through visual supervision. Together, they deliver the best overall performance in terms of harmonic mean (HM), highlighting their complementary effects.

#### A.2.7 Sensitive Analysis of Textual Prompt Template

The model’s performance is not highly sensitive to the choice of textual prompt templates. Following prior works such as CLIP and PromptFL, we use “a photo of a [CLASS]” as the default template. Since our method mainly optimizes the learnable prompt tokens ( $T_P$ ) while keeping the text tokens ( $T_T$ ) fixed, the initial template primarily serves as a starting point for semantic alignment and has limited impact on the final performance. We also conducted a small-scale experiment comparing different templates (e.g., “a picture of a [CLASS]”, “an image showing [CLASS]”), and the variation in Top-1 accuracy was less than 1%, indicating that the GPR-NIAM framework is robust to the choice of templates.

Table 13: Prompt template comparison on Caltech101, DTD, and UCF101. Best numbers in **bold**.

Template	Caltech101			DTD			UCF101		
	Base	Novel	HM	Base	Novel	HM	Base	Novel	HM
A photo of a [CLASS]	91.59	94.22	<b>92.89</b>	68.40	53.50	60.04	78.57	<b>70.28</b>	74.19
A picture of a [CLASS]	91.43	<b>94.34</b>	92.86	67.37	52.83	59.22	78.61	69.54	73.79
An image showing [CLASS]	<b>91.52</b>	93.88	92.68	<b>68.57</b>	<b>54.12</b>	<b>60.49</b>	<b>78.77</b>	70.14	<b>74.20</b>

## NeurIPS Paper Checklist

### 1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: This study outlines the research problem and its contributions in both the abstract and introduction.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

### 2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The limitations are discussed in the conclusion and future work sections.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [NA]

Justification: This paper does not involve theoretical assumptions.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

### 4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: This paper provides a detailed description of the experimental setup, including the range of parameter tuning, and the code will be uploaded as supplementary material.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

### 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The code will be uploaded as supplementary material.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [\[Yes\]](#)

Justification: The experimental setup is clearly described in both the main experiment section and the appendix.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

## 7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [\[Yes\]](#)

Justification: This paper adopts evaluation metrics and assessment methods consistent with existing studies.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer “Yes” if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The computational resources used are described in the experimental implementation details section.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

## 9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research conducted in this paper fully conforms to the NeurIPS Code of Ethics in every respect.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: As discussed in the introduction, the study substantially reduces communication costs by avoiding multiple rounds of exchange, which is advantageous for deployment under stringent latency or bandwidth limitations.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: The paper utilizes the fully open-source CLIP model and publicly available benchmark datasets that are widely used in existing literature. Therefore, no additional safeguards are necessary for responsible release.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

## 12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: the paper cites all relevant works, and all utilized assets are open-source and used in accordance with their licenses.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, [paperswithcode.com/datasets](https://paperswithcode.com/datasets) has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

## 13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The code developed for this work is included as anonymized supplementary material

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

## 14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]

Justification: The paper does not involve crowdsourcing nor research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

#### 15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA]

Justification: The paper does not involve crowdsourcing nor research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

#### 16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or originality of the research, declaration is not required.

Answer: [NA]

Justification: This paper only uses LLMs for grammar checking and editing.

Guidelines:

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.